

CPE 486/586: Deep Learning for Engineering Applications

07 Variational AutoEncoder

Spring 2026

Rahul Bhadani

Outline

1. Motivation
2. AutoEncoder
3. Variational Inference
4. Variational AutoEncoder

Motivation


0	1	0	1	0	1	1	0	0	1	0	0	0	0	0	0	1	0	1	
0	1	1	0	1	0	0	1	0	1	1	0	1	1	0	1	1	1	0	1
0	1	0	0	0	0	0	1	1	1	1	1	0	1	1	1	0	0	0	1

Estimating Probability Distributions for Generative Processes

Given $\mathbf{x}$, we aim to estimate $P(\mathbf{x})$ to provide a complete description of the probabilistic model that generates the observed data.

- ⚡ **Data Synthesis:** We want to learn generative processes capable of producing new datasets.
- ⚡ **Noise Reduction:** By discarding noise, we can more effectively learn underlying physical laws.
- ⚡ **Causality:** Generative models provide a framework to express and understand causal relations.

Example: *If we understand the generative process of wildfires, we can apply that knowledge across different regions, such as California or Australia.*




In this lecture, we model generative processes using Variational Autoencoders (VAEs).

AutoEncoder

0	1	1	0	0	0	1	1	1	1	1	1	1	0	1	1	1	0	0	0
0	1	0	1	1	0	0	0	0	1	1	1	1	1	0	0	1	1	0	1
1	1	0	0	0	1	1	1	1	0	0	1	0	1	0	0	0	1	1	1

What is AutoEncoder



Loss function is generally reconstruction loss:

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n ||\mathbf{x}_i - \mathbf{x}_i^d||_2^2$$

An **autoencoder** is a neural network, whose main job is to encode the input into a compressed representation, and then decode it back such that the reconstructed input is similar as possible to the original one.

It has three parts:

- 1 **Encoder:** Learns the most relevant features and transfers data into a latent space.
- 2 **Latent representation:** a middle layer that represents feature in some low dimensional space.
- 3 **Decoder:** it takes the reduced features from latent space and lift to the original dimension.

As it turns out, AutoEncoder is a neural network with a bottleneck.


Denoising Autoencoder

An AutoEncoder can end up learning an identity function (also called null function) where input=output.

To remedy, we add some noise to the input to corrupt it before feeding to AutoEncoder. Such an AE is called denoising AutoEncoder (DAE).

There are two ways to add noise to DAE:

- 1 Add Gaussian noise to the input.
- 2 Randomly set some input nodes to zero.
- 3 Masking out a certain percentage of the input's pixels in the case of image data.



Stacked Autoencoders

The stacked autoencoder is a *hack* to obtain a deep model by training only shallow ones in a **layer-wise, greedy fashion**.

- ⚡ **Phase 1:** Train the outermost encoder and decoder as a single shallow AE.
- ⚡ **Phase 2:** Encode the dataset; use the result to train the next inner shallow AE.
- ⚡ **Phase 3:** Repeat until reaching the innermost layers.
- ⚡ **Result:** Stack the pre-trained layers into one deep architecture.

Visualizing the Greedy Training Process

1. Train Outermost (E_1, D_1)




2. Train Middle (E_2, D_2)



3. Train Innermost (E_3, D_3)



4. Final Deep Hacked Stack



Reading List

- 1 Vincent, P., Larochelle, H., Bengio, Y., & Manzagol, P. A. (2008, July). Extracting and composing robust features with denoising autoencoders. In Proceedings of the 25th international conference on Machine learning (pp. 1096-1103). <https://dl.acm.org/doi/pdf/10.1145/1390156.1390294>
- 2 Alain, Guillaume, and Yoshua Bengio. "What regularized auto-encoders learn from the data-generating distribution." The Journal of Machine Learning Research 15, no. 1 (2014): 3563-3593. <https://www.jmlr.org/papers/volume15/alain14a/alain14a.pdf>
- 3 Vincent, Pascal, Hugo Larochelle, Isabelle Lajoie, Yoshua Bengio, Pierre-Antoine Manzagol, and Léon Bottou. "Stacked denoising autoencoders: Learning useful representations in a deep network with a local denoising criterion." Journal of machine learning research 11, no. 12 (2010).
- 4 Zhai, J., Zhang, S., Chen, J. and He, Q., 2018, October. Autoencoder and its various variants. In 2018 IEEE international conference on systems, man, and cybernetics (SMC) (pp. 415-419). IEEE.
- 5 Li, Pengzhi, Yan Pei, and Jianqiang Li. "A comprehensive survey on design and application of autoencoder in deep learning." Applied Soft Computing 138 (2023): 110176.
- 6 Michelucci, Umberto. "An introduction to autoencoders." arXiv preprint arXiv:2201.03898 (2022).
- 7 Bank, Dor, Noam Koenigstein, and Raja Giryes. "Autoencoders." Machine learning for data science handbook: data mining and knowledge discovery handbook (2023): 353-374.

Variational Inference

1	0	0	1	0	0	0	1	1	1	1	1	0	0	0	0	1	0	1	0
1	1	1	0	0	0	1	0	1	1	1	1	1	0	1	0	0	1	0	1
0	0	0	0	0	1	0	1	0	1	0	1	1	1	0	1	1	0	1	0

Motivation

We are interested in **posterior distribution calculation**.

Recall: Bayesian Framework

Given the parameter θ for a data distribution, we have prior as $p(\theta)$, a likelihood, if we write dataset as $\mathcal{D}$, then the posterior distribution can be calculated as:

$$p(\theta|\mathcal{D}) = \frac{p(\mathcal{D}|\theta)p(\theta)}{p(\mathcal{D})}$$

with $p(\mathcal{D}) = \int p(\mathcal{D}|\theta)p(\theta)d\theta$ as the *marginal distribution*.

θ denotes latent parameters that we want to infer.

So, what's the problem?

Motivation

We are interested in **posterior distribution calculation**.

Recall: Bayesian Framework

Given the parameter θ for a data distribution, we have prior as $p(\theta)$, a likelihood, if we write dataset as $\mathcal{D}$, then the posterior distribution can be calculated as:

$$p(\theta|\mathcal{D}) = \frac{p(\mathcal{D}|\theta)p(\theta)}{p(\mathcal{D})}$$

with $p(\mathcal{D}) = \int p(\mathcal{D}|\theta)p(\theta)d\theta$ as the *marginal distribution*.

θ denotes latent parameters that we want to infer.

So, what's the problem?

Computing $p(\mathcal{D})$ is **hard, intractable** due to the volume of data needed, and hence as a result, computing the posterior is hard too.

Variational Inference (VI) is one way to do so.

Variational Objective

We turn VI into an optimization problem: we choose a family of tractable distributions $\mathcal{Q}$ and find a member $q(\theta) \in \mathcal{Q}$ that is as close to $p(\theta|\mathcal{D})$ as possible.

What's the criteria for measuring the closeness between two distributions?

Variational Objective

We turn VI into an optimization problem: we choose a family of tractable distributions $\mathcal{Q}$ and find a member $q(\theta) \in \mathcal{Q}$ that is as close to $p(\theta|\mathcal{D})$ as possible.

What's the criteria for measuring the closeness between two distributions?

KL-Divergence, JS-divergence, etc.

Goal: Minimize KL divergence from $q(\theta)$ to $p(\theta|\mathcal{D})$

KL Divergence I

Definition

$$KL(q(\theta)||p(\theta|\mathcal{D})) = \int q(\theta) \log \frac{q(\theta)}{p(\theta|\mathcal{D})} d\theta$$

But KL divergence still depends on $p(\mathcal{D}) = \int p(\mathcal{D}|\theta)p(\theta)d\theta$.

To see, how, we can rewrite the KL integral using the definition of Expectation.

KL Divergence II

$$\begin{aligned}\text{KL}(q(\theta)||p(\theta|\mathcal{D})) &= \mathbb{E}_q \left[\log \frac{q(\theta)}{p(\theta|\mathcal{D})} \right] \\&= \mathbb{E}_q[\log q(\theta)] - \mathbb{E}_q[\log p(\theta|\mathcal{D})] \\&= \mathbb{E}_q[\log q(\theta)] - \mathbb{E}_q \left[\log \frac{p(\theta, \mathcal{D})}{p(\mathcal{D})} \right] \\&= \mathbb{E}_q[\log q(\theta)] - \mathbb{E}_q[\log p(\theta, \mathcal{D})] + \mathbb{E}_q[\log p(\mathcal{D})] \\&= \mathbb{E}_q[\log q(\theta)] - \mathbb{E}_q[\log p(\theta, \mathcal{D})] + \int q(\theta) \log p(\mathcal{D}) d\theta \\&= \mathbb{E}_q[\log q(\theta)] - \mathbb{E}_q[\log p(\theta, \mathcal{D})] + \log p(\mathcal{D})\end{aligned}$$

The last step is because $\log p(\mathcal{D})$ is a constant with respect to $q(\theta)$.

KL Divergence III

We can ignore $\log p(\mathcal{D})$ in our optimization problem.

Minimizing a quantity means maximizing its negative, and so we maximize the following quantity:

$$\begin{aligned}\text{ELBO}(q) &= -(\text{KL}(q(\theta) || p(\theta | \mathcal{D})) - \log p(\mathcal{D})) \\ &= -\left(\mathbb{E}_q[\log q(\theta)] - \mathbb{E}_q[\log p(\theta, \mathcal{D})] + \log p(\mathcal{D}) - \log p(\mathcal{D}) \right) \\ &= \mathbb{E}_q[\log p(\theta, \mathcal{D})] - \mathbb{E}_q[\log q(\theta)]\end{aligned}$$

We could expand joint probability and rewrite

KL Divergence IV

$$\begin{aligned}\text{ELBO}(q) &= \mathbb{E}_q[\log p(\mathcal{D}|\theta)] + \mathbb{E}_q[\log p(\theta)] - \mathbb{E}_q[\log q(\theta)] \\ &= \mathbb{E}_q[\log p(\mathcal{D}|\theta)] + \mathbb{E}_q\left[\log \frac{p(\theta)}{q(\theta)}\right] \\ &= \mathbb{E}_q[\log p(\mathcal{D}|\theta)] - \mathbb{E}_q\left[\log \frac{q(\theta)}{p(\theta)}\right] \\ &= \mathbb{E}_q[\log p(\mathcal{D}|\theta)] - \text{KL}(q(\theta)||p(\theta))\end{aligned}$$

We could also note that

$p(\mathcal{D}, \theta) = p(\mathcal{D})p(\theta, \mathcal{D})$ from the Bayes' rule.

KL Divergence V

Hence,

$$\text{ELBO}(q) = \mathbb{E}_q[\log p(\theta|\mathcal{D}) + \log p(\mathcal{D})] - \mathbb{E}_q[\log q(\theta)]$$

Taking constant out of the expectation, and distributing the expectation

$$\text{ELBO}(q) = \mathbb{E}_q[\log p(\theta|\mathcal{D})] - \mathbb{E}_q[\log q(\theta)] + \log p(\mathcal{D})$$

Since $\mathcal{D}$ is fixed, $\log p(\mathcal{D})$ is a constant with respect to q

$$\text{ELBO}(q) = \mathbb{E}_q[\log p(\theta|\mathcal{D})] - \mathbb{E}_q[\log q(\theta)] + \text{constant}$$

Variational Objective: The ELBO

The **Evidence Lower Bound (ELBO)** is defined as:

$$\text{ELBO}(q) = \mathbb{E}_q[\log p(\mathcal{D}|\theta)] - \text{KL}(q(\theta) \parallel p(\theta))$$

because the marginal is called *evidence* as well.

A $\log p(\mathcal{D})$ is constant wrt q , hence minimizing $\text{KL}(q(\theta) \parallel p(\theta|\mathcal{D}))$ is equivalent to maximizing ELBO. ELBO consists of two terms:

- **Expected log-likelihood:** Encourages q to place mass on parameters that explain the data well.
- **Negative KL to prior:** Encourages q to be close to the prior (regularizer).

VI maximizes a lower bound on the log marginal likelihood.

Question !

Show that $\log p(\mathcal{D}) \geq \text{ELBO}(q)$.

Question !

Show that $\log p(\mathcal{D}) \geq \text{ELBO}(q)$.

PROOF. —

$$\text{ELBO}(q) = -(\text{KL}(q(\theta) || p(\theta|\mathcal{D})) - \log p(\mathcal{D}))$$

$$\log p(\mathcal{D}) = \text{ELBO}(q) + (\text{KL}(q(\theta) || p(\theta|\mathcal{D})))$$

$$\log p(\mathcal{D}) \geq \text{ELBO}(q)$$

□

Maximizing ELBO finds approx distribution $q(\theta)$ for latent parameters θ that allows data to be predicted really well, but penalty occurs when $q(\theta)$ is far away from the prior $p(\theta)$.

What is Variational in VI?

We are not looking to find a single best value that maximizes the log likelihood, but rather a single best function.

Moving Beyond MLE

- ⚡ We are **not** looking to find a single best value that maximizes the log-likelihood.
- ⚡ Instead, we seek to identify the single **best function** (distribution).

The Methodology

- ⚡ To achieve this, we employ **variational calculus**.
- ⚡ The term "Variational" is derived directly from this mathematical framework.

Goal: Approximate a complex posterior by finding the closest member of a simpler family of distributions.

Variational Calculus, Briefly

Variational calculus (Calc is concerned with finding the functions that minimize or maximize quantities expressed as integrals. [**Pattern Recognition and Machine Learning (Christopher M. Bishop), Section 10.1**]

Unlike basic calculus, which focuses on derivatives of functions, variational calculus focuses on finding the function itself that produces the optimal outcome.

In Variational Calculus,

Find a function $y(x)$ that minimizes the functional $J[y] = \int F(x, y, y')dx$.
Here, $J[y]$ is not just a number but a **functional** — a mapping from functions to numbers.

The integral might represent things like total energy, time, or distance and our goal is to minimize it.

Some examples where Variational Calculus are useful

- 1 What is the shortest or smoothest path to the goal in case of autonomous navigation?
- 2 How do navigate while consuming the least energy?
- 3 How to avoid collision while being efficient?

The Euler-Lagrange Equation

The main tool in variational calculus is the Euler-Lagrange equation

$$\frac{d}{dx} \frac{\partial F}{\partial y'} - \frac{\partial F}{\partial y} = 0$$

which tells us the condition a function $y(x)$ must satisfy to minimize/maximize a given functional. In robotics, $y(x)$ could be robot's trajectory. F could be energy, cost, control effort, distance, etc.

Mean-field Variational Family

The next question, how do we choose Q . A common choice for the variational family Q is the **mean-field family**. Other options are *full-covariance Gaussians*, and *Normalizing Flows*.

The mean-field variational family assumes that the latent variables are independent of each other in the approximate posterior. Mathematically, if we have parameters $\theta = (\theta_1, \theta_2, \dots, \theta_M)$, the mean-field approximation factorizes as:

$$q(\theta) = \prod_{j=1}^M q_j(\theta_j)$$

Each q_j is a marginal variational distribution for the j -th parameter or group of parameters. The key point is that there are no cross-terms or dependencies between different θ_j 's in the approximation.

The term "mean-field" comes from physics, specifically from statistical mechanics. In physics, mean-field theory approximates a many-body problem by replacing interactions with an average or "mean" field. Similarly, in variational inference, we're approximating a complex posterior where variables are dependent by a simpler distribution where they're independent.

Mean-field Variational Inference with ELBO I

Recall: The ELBO Definition

$$\begin{aligned}\text{ELBO}(q) &= \mathbb{E}_q[\log p(\theta, \mathcal{D})] - \mathbb{E}_q[\log q(\theta)] \\ &= \int \prod_{i=1}^M q_i(\theta_i) \log p(\theta, \mathcal{D}) d\theta - \int \prod_{i=1}^M q_i(\theta_i)\end{aligned}$$

We optimize ELBO with respect to a single variational density $q_j(\theta_j)$ and keep others fixed.

Mean-field Variational Inference with ELBO II

Derivation for q_j

$$\begin{aligned}\text{ELBO}(q_j) &= \int \prod_{i=j}^M q_j(\theta_j) \log p(\theta, \mathcal{D}) d\theta - \int \prod_{i=j}^M q_j(\theta_j) \log \prod_{i=j}^M q_j(\theta_j) d\theta \\ &= \int \prod_{i=j}^M q_j(\theta_j) \log p(\theta, \mathcal{D}) d\theta - \int q_j(\theta_j) \log q_j(\theta_j) d\theta_j - \underbrace{\int \prod_{i \neq j}^M q_j(\theta_j) \log \prod_{i \neq j}^M q_j(\theta_j) d\theta_{-j}}_{\text{Fixed with respect to } q_j \theta_j} \\ &\propto \int \prod_{i=j}^M q_j(\theta_j) \log p(\theta, \mathcal{D}) d\theta - \int q_j(\theta_j) \log q_j(\theta_j) d\theta_j \\ &= \int q_j(\theta_j) \left(\int \prod_{i \neq j}^M q_j(\theta_j) \log p(\theta, \mathcal{D}) d\theta_{-j} \right) d\theta_j - \int q_j(\theta_j) \log q_j(\theta_j) d\theta_j \\ &= q_j(\theta_j) \mathbb{E}_{q(\theta_{-j})} [\log p(\theta, \mathcal{D})] d\theta_j - \int q_j(\theta_j) \log q_j(\theta_j) d\theta_j\end{aligned}$$

Mean-field Variational Inference with ELBO III

We could use variational calculus to find the optimal variational density $q_j^*(\theta_j)$.

We can further define $\log \tilde{p}(\mathcal{D}, \theta_j) = \mathbb{E}_{q(\theta_{-j})}[\log p(\theta, \mathcal{D})]d\theta_j + \text{Constant}$.

The constant is from the previous step just before the proportional.

Hence,

Relation to KL Divergence

$$\text{ELBO}(q_j) \propto \int q_j(\theta_j) \log \tilde{p}(\mathcal{D}, \theta_j) d\theta_j - \int q_j(\theta_j) \log q_j(\theta_j) d\theta_j$$

Merge the integral due to common integration variable

$$= \int q_j(\theta_j) \log \frac{\tilde{p}(\mathcal{D}, \theta_j)}{q_j(\theta_j)} d\theta_j$$

$$= - \int q_j(\theta_j) \log \frac{q_j(\theta_j)}{\tilde{p}(\mathcal{D}, \theta_j)} d\theta_j$$

$$= -\text{KL}(q_j(\theta_j) || \tilde{p}(\mathcal{D}, \theta_j)) \quad (\text{Using the definition of KL Divergence})$$

The ELBO is equal to the negative KL divergence up to an additive constant.

Hence, maximizing the ELBO is equivalent to minimizing the KL Divergence between $q_j(\theta_j)$ and $\tilde{p}(\mathcal{D}, \theta_j)$.

Mean-field Variational Inference with ELBO IV

Hence, under the mean-field assumption, the optimal variational density is given by

The Optimal Mean-Field Update

$$\begin{aligned} q_j^* &= \exp\left(\mathbb{E}_{q(\theta_{-j})}[\log p(\theta, \mathcal{D})]d\theta_j + \text{Constant}\right) \\ &= \frac{\mathbb{E}_{q(\theta_{-j})}[\log p(\theta, \mathcal{D})]d\theta_j}{\int \mathbb{E}_{q(\theta_{-j})}[\log p(\theta, \mathcal{D})]d\theta_j} \end{aligned}$$

Note that this is not an explicit solution because each optimal variational density depends on all others. Hence, the true solution will be iterative ones.

This process is called **Coordinated Ascent Variational Inference**.

Stochastic Gradients for Variational Inference

We can use gradient descent based methods for optimizing ELBO.

We need to compute gradients of the ELBO wrt to the parameters of q .

$$\mathcal{L}(\lambda) = \text{ELBO}(q_\lambda) = \mathbb{E}_{q_\lambda(\theta)}[\log p(\mathcal{D}|\theta)] - \text{KL}(q_\lambda(\theta) \parallel p(\theta))$$

where λ is the variational parameters such as mean and log-std of a Gaussian.

The KL term can often be computed analytically (e.g., when both q and p are Gaussians). The expected log-likelihood term, however, usually requires approximation.

In PyTorch, we can use reparameterization trick to estimate the gradients of the expectation.

Example: KL Divergence between Two Gaussians I

Derive the closed-form $KL(q\|p)$ where:

$$q(z) = \mathcal{N}(z; m, s^2) = \frac{1}{\sqrt{2\pi s^2}} \exp\left(-\frac{(z-m)^2}{2s^2}\right)$$

$$p(z) = \mathcal{N}(z; \mu_0, \sigma_0^2) = \frac{1}{\sqrt{2\pi\sigma_0^2}} \exp\left(-\frac{(z-\mu_0)^2}{2\sigma_0^2}\right)$$

Solution

$$\begin{aligned} KL(q\|p) &= \int_{-\infty}^{\infty} q(z) \log \frac{q(z)}{p(z)} dz \\ &= \int_{-\infty}^{\infty} q(z) \log q(z) dz - \int_{-\infty}^{\infty} q(z) \log p(z) dz \end{aligned}$$

Example: KL Divergence between Two Gaussians II

Taking the log of $q(z)$:

$$\begin{aligned}\log q(z) &= \log \left[\frac{1}{\sqrt{2\pi s^2}} \exp\left(-\frac{(z-m)^2}{2s^2}\right) \right] \\ &= -\frac{1}{2} \log(2\pi) - \frac{1}{2} \log(s^2) - \frac{(z-m)^2}{2s^2}\end{aligned}$$

Taking the log of $p(z)$:

$$\begin{aligned}\log p(z) &= \log \left[\frac{1}{\sqrt{2\pi\sigma_0^2}} \exp\left(-\frac{(z-\mu_0)^2}{2\sigma_0^2}\right) \right] \\ &= -\frac{1}{2} \log(2\pi) - \frac{1}{2} \log(\sigma_0^2) - \frac{(z-\mu_0)^2}{2\sigma_0^2}\end{aligned}$$

Example: KL Divergence between Two Gaussians III

Taking the difference:

$$\begin{aligned}\log q(z) - \log p(z) &= \left[-\frac{1}{2} \log(2\pi) - \frac{1}{2} \log(s^2) - \frac{(z-m)^2}{2s^2} \right] \\ &\quad - \left[-\frac{1}{2} \log(2\pi) - \frac{1}{2} \log(\sigma_0^2) - \frac{(z-\mu_0)^2}{2\sigma_0^2} \right] \\ &= \frac{1}{2} \log(\sigma_0^2) - \frac{1}{2} \log(s^2) - \frac{(z-m)^2}{2s^2} + \frac{(z-\mu_0)^2}{2\sigma_0^2} \\ &= \log \frac{\sigma_0}{s} - \frac{(z-m)^2}{2s^2} + \frac{(z-\mu_0)^2}{2\sigma_0^2}\end{aligned}$$

Example: KL Divergence between Two Gaussians IV

$$\begin{aligned} KL &= \mathbb{E}_q[\log q(z) - \log p(z)] \\ &= \mathbb{E}_q \left[\log \frac{\sigma_0}{s} - \frac{(z - m)^2}{2s^2} + \frac{(z - \mu_0)^2}{2\sigma_0^2} \right] \\ &= \log \frac{\sigma_0}{s} - \frac{1}{2s^2} \mathbb{E}_q[(z - m)^2] + \frac{1}{2\sigma_0^2} \mathbb{E}_q[(z - \mu_0)^2] \end{aligned}$$

Example: KL Divergence between Two Gaussians V

This is the variance of q :

$$\mathbb{E}_q \left[(z - m)^2 \right] = s^2$$

Add and subtract m inside the square:

$$\begin{aligned} \mathbb{E}_q \left[(z - \mu_0)^2 \right] &= \mathbb{E}_q \left[(z - m + m - \mu_0)^2 \right] \\ &= \mathbb{E}_q \left[(z - m)^2 \right] + 2(m - \mu_0) \mathbb{E}_q [z - m] + (m - \mu_0)^2 \\ &= s^2 + 2(m - \mu_0) \cdot 0 + (m - \mu_0)^2 \\ &= s^2 + (m - \mu_0)^2 \end{aligned}$$

where $\mathbb{E}_q [z - m] = 0$ since m is the mean of q .

Example: KL Divergence between Two Gaussians VI

$$\begin{aligned} KL &= \log \frac{\sigma_0}{s} - \frac{1}{2s^2} \cdot s^2 + \frac{1}{2\sigma_0^2} \left[s^2 + (m - \mu_0)^2 \right] \\ &= \log \frac{\sigma_0}{s} - \frac{1}{2} + \frac{s^2 + (m - \mu_0)^2}{2\sigma_0^2} \\ &= \log \frac{\sigma_0}{s} + \frac{s^2 + (m - \mu_0)^2}{2\sigma_0^2} - \frac{1}{2} \end{aligned}$$

Example: KL Divergence between Two Gaussians VII

Substituting $\mu_0 = 0$ and $\sigma_0 = 1$:

$$\begin{aligned} KL &= \log \frac{1}{s} + \frac{s^2 + m^2}{2 \cdot 1} - \frac{1}{2} \\ &= -\log s + \frac{s^2 + m^2}{2} - \frac{1}{2} \end{aligned}$$

References

- ⚡ Section 10.1 of Pattern Recognition and Machine Learning (Christopher M. Bishop 2006)
- ⚡ Stochastic Variational Inference <https://arxiv.org/pdf/1206.7051>

Variational AutoEncoder

0	0	0	0	0	1	1	0	1	1	1	0	1	0	1	1	0	0	1	1
0	0	1	0	1	0	1	1	1	1	0	1	1	1	0	1	1	0	1	1
1	1	1	0	0	0	1	0	0	0	0	0	0	1	0	0	1	1	1	0